

✓ We are building models and forecast 30 days after for Continental Breakfast

We should first check the data and the trend:

```
import warnings
with warnings.catch_warnings():
    warnings.simplefilter("ignore")
import pandas as pd
import statsmodels.api as sm
import plotly.express as px
```

Import data, read the data and see the data descriptions.

```
df = pd.read_excel("IMB465-XLS-ENG.xls")
df = df.rename(columns={'Continental B/F': 'Continental_BF', 'North Indian B/F': 'North_Indian_BF'})
print(df.info()) #information of the dataset
print(df.describe().round(2))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 115 entries, 0 to 114
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  115 non-null   datetime64[ns]
1   BKFST_OCCUP           115 non-null   int64
2   Idly                  115 non-null   int64
3   Dosa                  115 non-null   int64
4   Chutney               115 non-null   int64
5   Sambar                115 non-null   int64
6   Continental_BF        115 non-null   int64
7   North_Indian_BF      115 non-null   int64
8   Omellette             115 non-null   int64
dtypes: datetime64[ns](1), int64(8)
memory usage: 8.2 KB
None
```

	BKFST_OCCUP	Idly	Dosa	Chutney	Sambar	Continental_BF \
count	115.00	115.00	115.00	115.00	115.00	115.00
mean	215.16	59.27	25.48	131.34	131.47	41.30
std	19.38	7.03	11.38	10.87	11.12	5.79
min	151.00	40.00	0.00	95.00	95.00	25.00
25%	200.00	54.00	18.00	125.00	125.00	40.00
50%	215.00	61.00	22.00	130.00	130.00	40.00
75%	228.50	65.00	27.00	135.00	136.00	44.00
max	259.00	70.00	62.00	160.00	160.00	55.00


	North_Indian_BF	Omellette
count	115.00	115.00
mean	4.35	13.59
std	1.99	5.26
min	0.00	5.00
25%	2.50	10.00
50%	5.00	13.00
75%	6.00	20.00
max	8.00	22.00

Remodify the data with only the columns we need.

```
df = df[['Date', 'Continental_BF']]
df.index = df.Date
df.index.freq = 'D'
df = df.drop('Date', axis = 1)
```


Show the data trend

```
fig = px.line(x = df.index, y = df['Continental_BF'], labels = {'x': 'Date', 'y': "Continental_BF"})
fig.show()
```



Plot the seasonality charts to check the seasonality.

```
from statsmodels.tsa.seasonal import seasonal_decompose
decompose_data = seasonal_decompose(df, model="add")
decompose_data.plot();
```




✓ Moving Average

Moving average can show us the data trend more intuitively.

```
PERIODS_AHEAD = 30
periods = 7
rolling_MA = df.copy()
rolling_MA['MovingAverage'] = df.Continental_BF.rolling(periods).mean()
print(rolling_MA.head(14))
fig = px.line(rolling_MA, x = rolling_MA.index, y = ['Continental_BF', 'MovingAverage'],
              labels = {'value': 'Continental_BF', 'date': 'Date', 'variable': 'Variable'})
fig.show()
```

fig.show()

Date	Continental_BF	MovingAverage
2012-10-01	25	NaN
2012-10-02	25	NaN
2012-10-03	25	NaN
2012-10-04	35	NaN
2012-10-05	41	NaN
2012-10-06	30	NaN
2012-10-07	40	31.571429
2012-10-08	40	33.714286
2012-10-09	40	35.857143
2012-10-10	40	38.000000
2012-10-11	40	38.714286
2012-10-12	40	38.571429
2012-10-13	44	40.571429
2012-10-14	49	41.857143



✓ The Forecasting Models

Now we start to build models.

```
from statsmodels.tsa.holtwinters import SimpleExpSmoothing
from statsmodels.tsa.holtwinters import ExponentialSmoothing
```


✓ ETS models

For the first 3 models, we use exponential smoothing models (sample, double-holt and triple-winter). We just use `optimized = True` to get the optimized solution for the models.

✓ Sample

```
ets_model = SimpleExpSmoothing(df)
ets_fit = ets_model.fit(optimized=True)
# print(ets_fit.fittedvalues)


PERIODS_AHEAD = 30
fig = px.line(x = df.index, y = df['Continental_BF'], labels = {'x': 'Date', 'y': 'Continental_BF'})
fig.add_scatter(x = ets_fit.fittedvalues.index, y = ets_fit.fittedvalues, name = 'Fitted Values')
fig.add_scatter(x = ets_fit.forecast(PERIODS_AHEAD).index, y = ets_fit.forecast(PERIODS_AHEAD), name = 'Forecast')
fig.show()
```



▼ Holt

```
ets_model = ExponentialSmoothing(df, trend='add',
                                seasonal=None, initialization_method = 'estimated')
ets_fit = ets_model.fit(optimized = True)
forecast = ets_fit.forecast(PERIODS_AHEAD)
# print(forecast)

fig = px.line(x = df.index, y = df['Continental_BF'], labels = {'x': 'Date', 'y': 'Demand'})
fig.add_scatter(x = ets_fit.fittedvalues.index, y = ets_fit.fittedvalues, name = 'Fitted Values')
fig.add_scatter(x = ets_fit.forecast(PERIODS_AHEAD).index, y = ets_fit.forecast(PERIODS_AHEAD), name = 'Forecast')
fig.show()
```




▼ Winter

```

ets_model = ExponentialSmoothing(df, trend='add', seasonal='mul', seasonal_periods=7)
ets_fit = ets_model.fit(optimized=True)
forecast = ets_fit.forecast(PERIODS_AHEAD)
# print(ets_fit.forecast(PERIODS_AHEAD))

fig = px.line(x = df.index, y = df['Continental_BF'], labels = {'x': 'Date', 'y': 'Demand'})
fig.add_scatter(x = ets_fit.fittedvalues.index, y = ets_fit.fittedvalues, name = 'Fitted Values')
fig.add_scatter(x = ets_fit.forecast(PERIODS_AHEAD).index, y = ets_fit.forecast(PERIODS_AHEAD), name = 'Forecast')
fig.show()

```



ARIMA

✓ Check for stationarity

```

from statsmodels.tsa.stattools import adfuller
adf = adfuller(df)

print(f'ADF Statistic {adf[0]}') # adf[0] - returns the ADF statistic value
print(f'p-value {adf[1]}') # adf[1] - returns the p-value -- if this value is high, the data is non-stationary

```

```

ADF Statistic -6.140300480128494
p-value 8.001298246056169e-08

```


✓ Explore autocorrelation

The plots help us determine what the parameter p and q should be.

```

from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
plot_acf(df).show() #q=1
plot_pacf(df).show() #p=1 or 2

```



✓ Search for the best ARIMA model using pmdarima

pmdarima automates the search for the best ARIMA model, it goes through many p,q,d combinations, and selects the model with the lowest AIC

```
#install pmdarima package
!pip install pmdarima
```

```
Requirement already satisfied: pmdarima in /usr/local/lib/python3.10/dist-packages (2.0.4)
Requirement already satisfied: joblib<=0.11 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.3.2)
Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.2.3)
Requirement already satisfied: numpy<=1.21.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.23.5)
Requirement already satisfied: pandas<=0.19 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.5.3)
Requirement already satisfied: scikit-learn<=0.22 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.2.2)
Requirement already satisfied: scipy<=1.3.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.11.4)
Requirement already satisfied: statsmodels<=0.13.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (0.14.1)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (2.0.7)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (67.7.0)
Requirement already satisfied: packaging<=17.1 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (23.2)
Requirement already satisfied: python-dateutil<=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas<=0.19->pmdarima) (2.8.1)
Requirement already satisfied: pytz<=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas<=0.19->pmdarima) (2020.1)
Requirement already satisfied: threadpoolctl<=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn<=0.22->pmdarima) (2.0.0)
Requirement already satisfied: patsy<=0.5.4 in /usr/local/lib/python3.10/dist-packages (from statsmodels<=0.13.2->pmdarima) (0.5.4)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy<=0.5.4->statsmodels<=0.13.2->pmdarima) (1.16.0)
```

```
import pmdarima as pm

model_fit = pm.auto_arima(df,
                          start_p=1, start_q=1, start_d=1,
                          max_p=12, max_q=12, max_d=12,
                          seasonal=False, #since we already observed no seasonality in the data
                          error_action='ignore',
                          suppress_warnings=True,
                          trace=True,
                          stepwise=True,
                          stationary=True)

print(model_fit.summary())
```

```
Performing stepwise search to minimize aic
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=692.761, Time=0.14 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=733.400, Time=0.03 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=695.003, Time=0.09 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=694.928, Time=0.06 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=1186.400, Time=0.01 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=694.217, Time=0.86 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=694.750, Time=0.39 sec
ARIMA(0,0,2)(0,0,0)[0] intercept : AIC=693.383, Time=0.17 sec
ARIMA(2,0,0)(0,0,0)[0] intercept : AIC=692.523, Time=0.20 sec
ARIMA(3,0,0)(0,0,0)[0] intercept : AIC=694.454, Time=0.66 sec
ARIMA(3,0,1)(0,0,0)[0] intercept : AIC=696.514, Time=0.69 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=inf, Time=0.12 sec
```

```
Best model: ARIMA(2,0,0)(0,0,0)[0] intercept
Total fit time: 3.445 seconds
```

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:      115
Model:                SARIMAX(2, 0, 0)      Log Likelihood      -342.262
Date:                Sun, 04 Feb 2024      AIC              692.523
Time:                04:46:03      BIC              703.503
Sample:              10-01-2012      HQIC             696.980
                  - 01-23-2013
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
intercept	22.2989	2.846	7.836	0.000	16.722	27.876
ar.L1	0.6629	0.090	7.340	0.000	0.486	0.840
ar.L2	-0.2035	0.095	-2.152	0.031	-0.389	-0.018
sigma2	22.4350	2.544	8.817	0.000	17.448	27.422

```
=====
Ljung-Box (L1) (Q):      0.08      Jarque-Bera (JB):      2.59
Prob(Q):                0.77      Prob(JB):              0.27
Heteroskedasticity (H):  1.00      Skew:                  0.05
Prob(H) (two-sided):    1.00      Kurtosis:              3.73
=====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

✓ Result Validating (5 models)

```
# Let us build a function that calculates MAPE to automate some steps further:
def calculate_mape(actual_values, predicted_values):
    return ((abs((actual_values - predicted_values) / actual_values))*100).mean()
```

We decide to split the train and test data by 0.9 : 0.1 (so we include new year days in our train data), since we observed **a big surge during New Year days, and this might affect our validation.**

```
# Create a train and test sets:
train, test = df.iloc[:round(len(df) * 0.9)], df.iloc[round(len(df) * 0.9):]
print(len(test))
print(train.index) # dates in train data set
```

```
11
DatetimeIndex(['2012-10-01', '2012-10-02', '2012-10-03', '2012-10-04',
               '2012-10-05', '2012-10-06', '2012-10-07', '2012-10-08',
               '2012-10-09', '2012-10-10',
               ...,
               '2013-01-03', '2013-01-04', '2013-01-05', '2013-01-06',
```

```
'2013-01-07', '2013-01-08', '2013-01-09', '2013-01-10',
'2013-01-11', '2013-01-12'],
dtype='datetime64[ns]', name='Date', length=104, freq='D')
```

✓ Cross-validate the results (ARIMA)

```
arima = pm.auto_arima(train,
                      start_p=1, start_q=1, start_d=1,
                      max_p=12, max_q=12, max_d=12,
                      seasonal=False,
                      error_action='ignore',
                      suppress_warnings=True,
                      trace = True,
                      stepwise = True,
                      stationary=False)

Performing stepwise search to minimize aic
ARIMA(1,0,1)(0,0,0)[0] : AIC=654.446, Time=0.26 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=1072.766, Time=0.04 sec
ARIMA(1,0,0)(0,0,0)[0] : AIC=653.901, Time=0.07 sec
ARIMA(0,0,1)(0,0,0)[0] : AIC=949.376, Time=0.15 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=inf, Time=0.15 sec
ARIMA(2,0,1)(0,0,0)[0] : AIC=657.430, Time=0.24 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=630.095, Time=0.13 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=667.106, Time=0.01 sec
ARIMA(2,0,0)(0,0,0)[0] intercept : AIC=628.661, Time=0.16 sec
ARIMA(3,0,0)(0,0,0)[0] intercept : AIC=630.657, Time=0.54 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=630.313, Time=0.30 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=628.456, Time=0.35 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=630.386, Time=0.11 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=630.417, Time=0.62 sec
ARIMA(0,0,2)(0,0,0)[0] intercept : AIC=629.475, Time=0.17 sec
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=632.276, Time=0.38 sec

Best model: ARIMA(1,0,1)(0,0,0)[0] intercept
Total fit time: 3.724 seconds
```

The Validation of all models

✓ Moving Aaverage

```
# Step 1: Calculate the rolling moving average for the training set
rolling_MA_train = train.copy()
rolling_MA_train['MovingAverage'] = train['Continental_BF'].rolling(periods).mean()

# Step 2: Extrapolate the trend for the test set based on the last value of the rolling moving average in the training set
last_ma_value_train = rolling_MA_train['MovingAverage'].iloc[-1]

# Create a DataFrame to store the forecasted values for the test set
forecast_test = pd.DataFrame(index=test.index)

# Fill the forecast DataFrame with the last value of the moving average in the training set
forecast_test['Forecast'] = last_ma_value_train

# Calculate the Mean Absolute Percentage Error (MAPE)
actual_values = test['Continental_BF']
forecast_values = forecast_test['Forecast']
mape = (abs(actual_values - forecast_values) / actual_values).mean() * 100

# Print the MAPE
print("Mean Absolute Percentage Error (MAPE): {:.2f}%".format(mape))
forecast_ma=forecast_values

Mean Absolute Percentage Error (MAPE): 8.39%
```

✓ ETS and ARIMA MAPES


```

#Simple ETS
ets_model_ets = SimpleExpSmoothing(train).fit(optimized=True)
forecast_ets = ets_model_ets.forecast(len(test))
print("MAPE for the Simple ETS model:", calculate_mape(test["Continental_BF"], forecast_ets), "%")

# Holt
# Create forecast using Holt's model here

ets_model_ets_H = ExponentialSmoothing(train, trend='add',seasonal=None, initialization_method = 'estimated').fit(optimized = Tr
forecast_ets_H = ets_model_ets_H.forecast(len(test))
print("MAPE for the Holt model:", calculate_mape(test["Continental_BF"], forecast_ets_H), "%")

# Winters
# Create forecast using Winters' model here

ets_model_ets_W = ExponentialSmoothing(train, trend='add', seasonal='mul', seasonal_periods=7).fit(optimized = True)
forecast_ets_W = ets_model_ets_W.forecast(len(test))
print("MAPE for the Winter model:", calculate_mape(test["Continental_BF"], forecast_ets_W), "%")

#ARIMA
print('MAPE for the ARIMA model:', calculate_mape(test['Continental_BF'], arima.predict(len(test))), "%")

    MAPE for the Simple ETS model: 8.715739790724 %
    MAPE for the Holt model: 8.191732186507025 %
    MAPE for the Winter model: 11.429565900517401 %
    MAPE for the ARIMA model: 8.211194365546056 %

# Plot the results (fitted values and forecasts) using all 4 models (ETS, Holt's, Winters, ARIMA)

fig = px.line(x = df.index, y = df['Continental_BF'], labels = {'x': 'Date', 'y': 'Continental_BF'})

fig.add_scatter(x = ets_model_ets.fittedvalues.index, y = ets_model_ets.fittedvalues, name = 'ETS Fitted Values',
    line=dict(color='red') )

fig.add_scatter(x = test.index, y = forecast_values, name = 'Moving Average Forecast',
    line=dict(color='yellow', dash='dash') )

fig.add_scatter(x = test.index, y = ets_model_ets.forecast(len(test)), name = 'ETS Forecast',
    line=dict(color='red', dash='dash') )


# fig.add_scatter(x = ets_model_ets_H.fittedvalues.index, y = ets_model_ets_H.fittedvalues, name = 'ETS Holt Fitted Values',
#     line=dict(color='blue') )
fig.add_scatter(x = test.index, y = ets_model_ets_H.forecast(len(test)), name = 'ETS Holt Forecast',
    line=dict(color='blue', dash='dash') )

# fig.add_scatter(x = ets_model_ets_W.fittedvalues.index, y = ets_model_ets_W.fittedvalues, name = 'ETS Winter Fitted Values',
#     line=dict(color='black') )
fig.add_scatter(x = test.index, y = ets_model_ets_W.forecast(len(test)), name = 'ETS Winter Forecast',
    line=dict(color='black', dash='dash') )

fig.add_scatter(x = arima.predict(len(test)).index, y = arima.predict(len(test)), name = 'ARIMA Forecast',
    line=dict(color='green', dash='dash') )

fig.show()

```



MAPE for the Holt model: 0.08191732186507027

MAPE for the ARIMA model: 0.08211194365546058

These two models give us the lowest MAPE, indicating a better fit of testing data, they can be best models for future forecasting.

To decide which to use, we test for 100 future days:

```
extended_dates = pd.date_range(start=df.index.min(), periods=len(df) + 100, freq=df.index.freq)

# Plotting using Plotly Express
fig = px.line(x=df.index, y=df['Continental_BF'], labels={'x': 'Date', 'y': 'Continental_BF'})


fig.add_scatter(x=ets_model_ets.fittedvalues.index, y=ets_model_ets.fittedvalues, name='ETS Fitted Values',
                line=dict(color='red'))
fig.add_scatter(x=ets_model_ets.forecast(len(test) + 100).index, y=ets_model_ets.forecast(len(test) + 100), name='ETS Forecast',
                line=dict(color='red', dash='dash'))

# Add similar modifications for other forecasting models

fig.add_scatter(x=ets_model_ets_H.forecast(len(test) + 100).index, y=ets_model_ets_H.forecast(len(test) + 100), name='ETS Holt F',
                line=dict(color='blue', dash='dash'))

fig.add_scatter(x=ets_model_ets_W.forecast(len(test) + 100).index, y=ets_model_ets_W.forecast(len(test) + 100), name='ETS Winter',
                line=dict(color='black', dash='dash'))

fig.add_scatter(x=arima.predict(len(test) + 100).index, y=arima.predict(len(test) + 100), name='ARIMA Forecast',
                line=dict(color='green', dash='dash'))
```



Even Holt-ETS model has a lower MAPE than ARIMA model, we still choose ARIMA model, since the Holt-ETS model only use last several days' data, and makes the forecasting data increasing continuously, it will only provide accuracy for several days after (not as long as 30 days).

ARIMA model includes more data and smoothed the trend in the forecasting, and it can be less risky if we predict longer future data. So we decide to use ARIMA model as our final model.

✓ Forecasting by ARIMA and Save files

```
from google.colab import files
```

```
#forecast by ARIMA
```

```
arima2 = pm.auto_arima(df,
                        start_p=1, start_q=1, start_d=1,
                        max_p=12, max_q=12, max_d=12,
                        seasonal=True,
                        error_action='ignore',
                        suppress_warnings=True,
                        trace = True,
                        stepwise = True,
                        stationary=False)
```

```
forecast_arima = arima2.predict(30)
```

```
print(forecast_arima)
```

```
Performing stepwise search to minimize aic
```

```
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=692.761, Time=0.27 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=733.400, Time=0.05 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=695.003, Time=0.11 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=694.928, Time=0.15 sec
ARIMA(0,0,0)(0,0,0)[0] : AIC=1186.400, Time=0.04 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=694.217, Time=0.91 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=694.750, Time=0.58 sec
ARIMA(0,0,2)(0,0,0)[0] intercept : AIC=693.383, Time=0.17 sec
ARIMA(2,0,0)(0,0,0)[0] intercept : AIC=692.523, Time=0.41 sec
ARIMA(3,0,0)(0,0,0)[0] intercept : AIC=694.454, Time=0.61 sec
ARIMA(3,0,1)(0,0,0)[0] intercept : AIC=696.514, Time=0.57 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=inf, Time=0.12 sec
```

```
Best model: ARIMA(2,0,0)(0,0,0)[0] intercept
```

```
Total fit time: 4.019 seconds
```

```
2013-01-24 42.718049
```

```
2013-01-25 41.664753
```

```
2013-01-26 41.227342
```

```
2013-01-27 41.151689
```

```
2013-01-28 41.190536
```

```
2013-01-29 41.231682
```

2013-01-30	41.251054
2013-01-31	41.255524
2013-02-01	41.254546
2013-02-02	41.252988
2013-02-03	41.252154
2013-02-04	41.251919
2013-02-05	41.251932